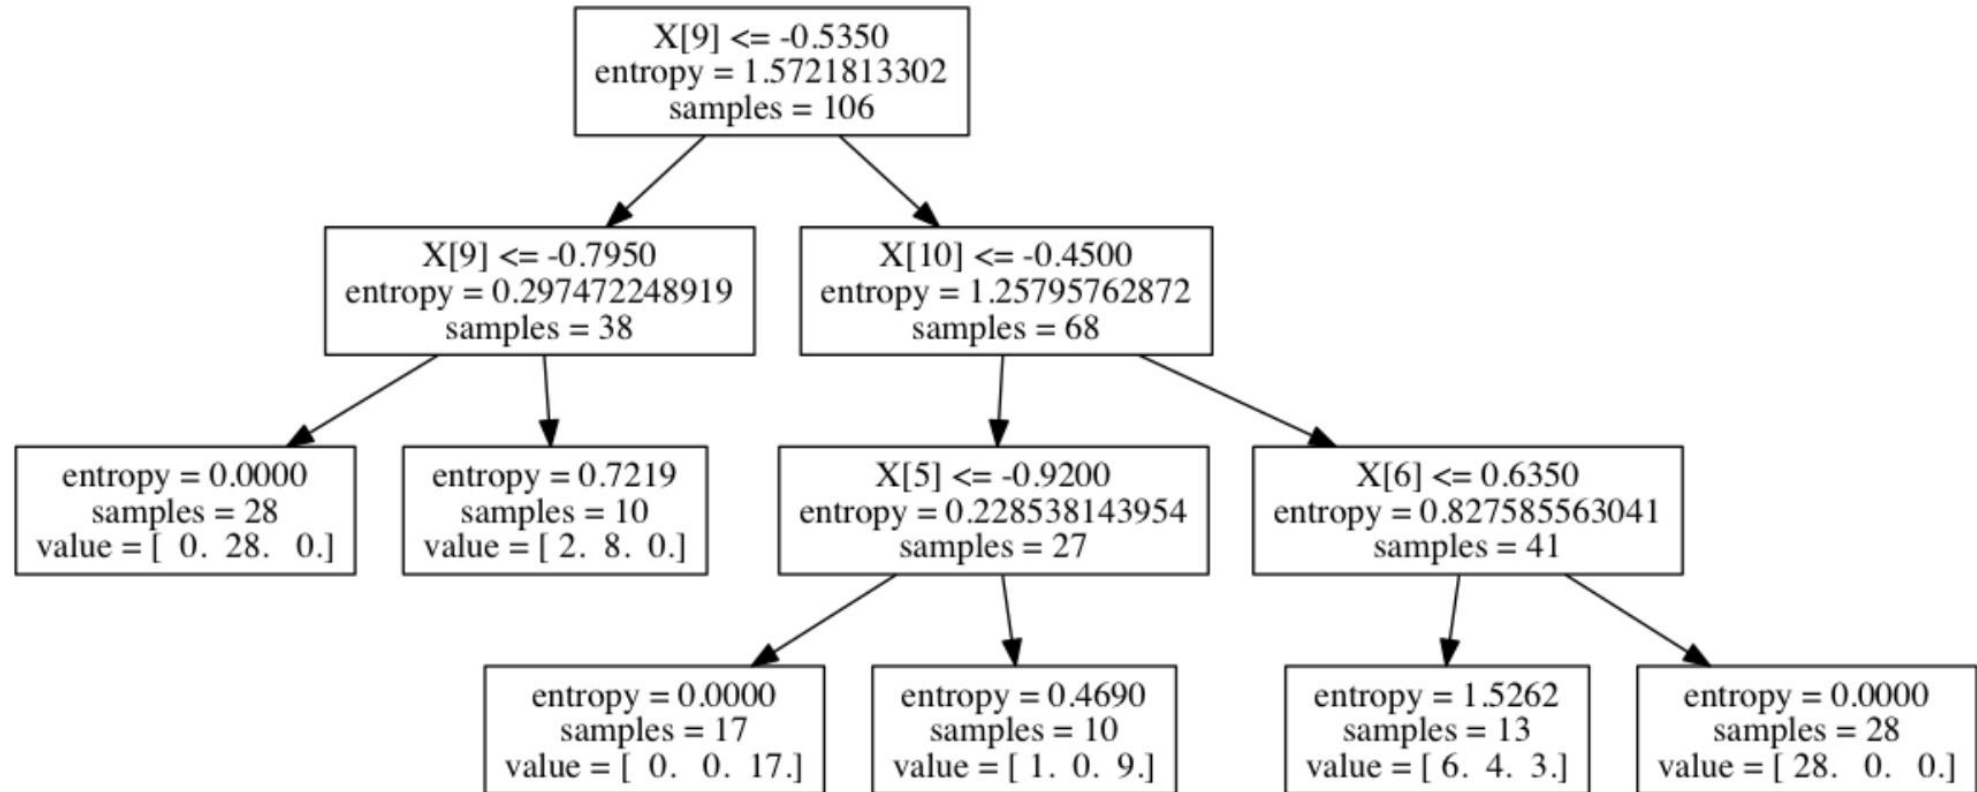
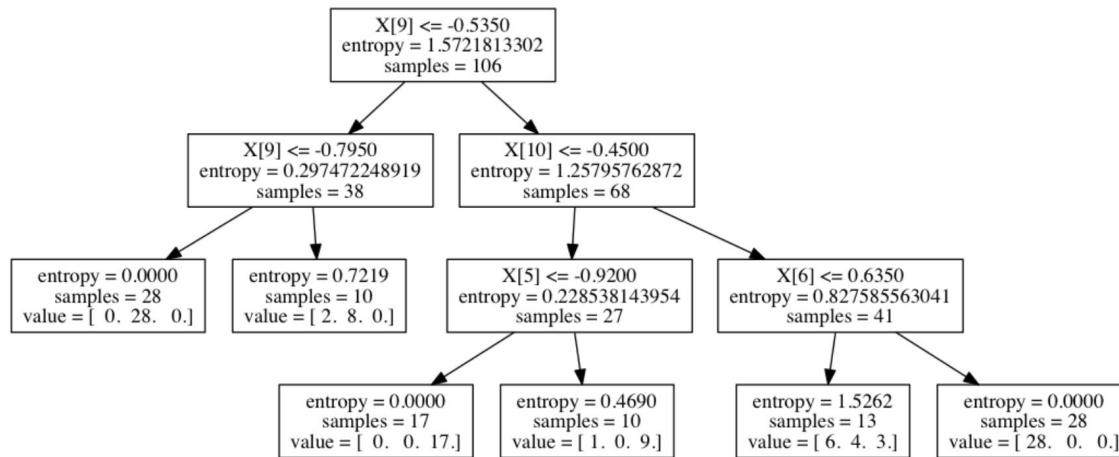


Decision trees

$$G(D, x) = \text{Entropy}(D) - \sum_{v \in x} \frac{|D_v|}{|D|} \text{Entropy}(D_v)$$



Decision trees



advantages:

- ease of interpretation
- handles continuous and discrete features
- invariant to monotone transformation of features
- variable selection automated
- **low bias** (deep trees)

disadvantages:

- **high variance**
- overfitting



Bias-variance tradeoff

- Error due to bias
- Error due to variance
- Irreducible error



Low bias

High bias

Low
variance




High
variance




Bias-variance tradeoff

- Error due to bias
- Error due to variance
- Irreducible error



bagging/boosting



Bagging


- Simulate the notion of different train set samples:
 1. sample data points from train set to create new train set
 2. fit low bias high variance model on new train set
 3. repeat steps (1) and (2) T times
 4. average the predictions of the T models

$$\hat{f}(x, \theta) = \frac{1}{T} \sum_{t=1}^T f_t(x, \theta)$$

Bagging: Random Forests

- Train set contains n data points with m features.
- Construct T low bias high variance decision trees by following these steps:
 1. Sample n data points at random **with replacement** from the train set.
 2. At each node, select $h \ll m$ features at random and compute the best split using only these h features.
 3. Each tree is grown to the largest extent possible. There is no pruning or early stopping.
- Step 3 ensures that the bagged models are low bias by learning deep complex decision trees.

Bagging: Random Forests



Boosting

- reduce the bias of a high bias low variance model
- turning an ensemble of weak learners into a strong learner
- the meta-model is additive, i.e. adaboost:

$$\hat{f}(x, \theta) = \sum_{t=1}^T \alpha_t f_t(x, \theta)$$

Boosting: adaboost

Initialize the weights $D_1(i) = 1/n, i = 1, 2, \dots, n$.

$$y_i \in \{-1, +1\}$$

For $t = 1$ to T :

1. Fit a weak classifier $f_t(x, \theta)$ to the trainset data using weights $D_1(i)$.


2. Set $\alpha_t = \frac{1}{2} \ln\left(\frac{1-\text{error}}{\text{error}}\right)$.

3. Update weights:


$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i f_t(x_i, \theta))}{Z_t},$$

where Z_t is a normalizing factor that makes sure that $\sum D_{t+1}(i) = 1$.

$$\hat{f}(x, \theta) = \sum_{t=1}^T \alpha_t f_t(x, \theta)$$




Boosting: adaboost



- The weight of a weak model in the boosted meta-model increases exponentially as the error approaches 0. Better models are given exponentially more weight.
- The weight is zero if the error rate is 0.5. A model with 50% accuracy is no better than random guessing, so it is ignored.
- The weight decreases exponentially as the error approaches 1. A negative weight is given to classifiers with worse than 50% accuracy. “Whatever that classifier says, do the opposite!”.

$$\hat{f}(x, \theta) = \sum_{t=1}^T \alpha_t f_t(x, \theta)$$

Boosting: adaboost



Boosting: Gradient Boosting

$$\hat{f}(x, \theta) = \sum_{t=1}^T f_t(x, \theta)$$

1. Fit a model $f_1(x, \theta) = y$
2. Fit a model to the **residuals** $h_1(x) = y - f_1(x, \theta)$
3. Create a new model $f_2(x, \theta) = f_1(x, \theta) + h_1(x)$

$$f_0(x, \theta) = \frac{1}{n} \sum_{i=1}^n y_i$$

$$f_{t+1}(x, \theta) = f_t(x, \theta) + h_t(x)$$

